

192.151 · Introduction to Deep Learning · 2026S

Project 13: Text to Speech

Fine-tuning VITS into a custom voice

Dima · Emir · Ilya

TU Wien



The goal

Just running a pretrained TTS model isn't a project, it already speaks well. We wanted to **change the voice itself**.


- ▶ **Task:** map text → audio, then **fine-tune** the model onto a new voice, a Star Wars **B1 battle-droid**.
- ▶ **Why a droid?** the effect is **audible even without metrics**: you can *hear* that fine-tuning worked.





Project architecture: built to move fast

Three of us on one model, so I scaffolded the repo **contract-first**: each part talks only through small fixed interfaces, never each other's internals.



- → **swap the model without touching eval**, and all three of us build in **parallel**.
- **Offline mocks + smoke test**: the whole pipeline wires up in seconds, no GPU. Plus pinned deps · fixed seed · `ruff` · `loguru` for reproducibility.





Three ways to change a TTS voice

Approach	What	Verdict
Pre-process	change the input (e.g. phoneme tweaks for an accent)	limited
Post-process	filter the output audio after the model	sounds unnatural
Bake into the model	fine-tune on data in the target voice	the real path ✓

Pre-processing is where I wanted the **accent**; I tried **post-processing** for the droid and it sounded unnatural, so I **bake the droid voice into the model**: it produces the voice *natively* from text, not as an after-effect.





Where to start: the ideas

Running a pretrained TTS isn't a project. I wanted to **change the voice itself**, and brainstormed four directions:

- **B1 droid voice:** a timbre change, instantly audible
- **Heavy accent:** phoneme-level manipulation
- **Emotions:** prosody control
- **Stuttering:** rhythm manipulation

I dropped stuttering & emotions (too shallow) and kept the **droid** as the core, accent as future work.



Our choice: bake in a droid voice

- **Core:** fine-tune on a **B1 droid voice**; keep the **accent** as future work.
- Post-processing sounded bad, so instead I used **RVC** to build the training data.

Why this idea fits the project:

- **One consistent voice** across the whole dataset.
- A unique effect, **audible even without metrics**.
- Injecting a voice *into training*, a real way to play with the learning itself.





Building the dataset — RVC

I needed 24h of paired text+audio in the droid voice, can't record that. So I **built it with RVC** (voice conversion): it repaints the **timbre** but keeps the **pronunciation**, so the original transcripts stay valid.

LJSpeech (13,100 clips) → RVC → B1 droid → resample 22.05 kHz → 📁 Hub

- ▶ Published: `Dmi1tr13/ljspeech-b1`
- ▶ 13,100 paired clips, a full droid dataset, for free



scan → the dataset



The trade-off I chose

Because the data is made by RVC, the model learns from **RVC's output**, not a real droid recording.

- ↓ **Quality can't beat the teacher:** RVC is the upper bound.
- ↑ But the model makes the droid voice **itself**, end-to-end, **no second model in production**.

Post-processing = run RVC every time. Baking = the voice is built in. A conscious trade: a bit less ceiling, full autonomy. (Ilya measures exactly where that ceiling sits, later.)






Why VITS

- **End-to-end:** text $\rightarrow$ waveform in **one model**: no separate vocoder, no intermediate spectrogram to predict then re-synthesize.
- **Phoneme-level access:** needed for the accent idea (future work).
- **Strong pretrained base** to fine-tune from: `kakao-enterprise/vits-ljs`.

Traditional TTS, two stages



VITS = a conditional VAE with a normalizing-flow prior, a GAN decoder, and a stochastic duration predictor. (Kim, Kong & Son, ICML 2021)



The generator: why it isn't enough

In training the model sees **both** text and the real audio, and learns to rebuild the waveform:



- ▶ Trained on reconstruction (L1 on the mel) **alone**, the audio is **weak**: blurry, muffled, averaged.
- ▶ L1 can't capture fine high-frequency detail → we need a **discriminator**.

At inference there's no audio, the model never transcribes; it only goes **text** → **speech**. The real audio is used in training only, to teach the text-prior (the KL term).



Discriminator: built from scratch

Hugging Face ships VITS **inference-only**: no discriminator. We wrote it, HiFi-GAN style:




Without the discriminator the decoder never learns realistic audio, it's the missing half of the VITS objective.



Lining up text & audio: MAS

A few text tokens, hundreds of audio frames, and **no labels** for which frame belongs to which token. **Monotonic Alignment Search (MAS)** finds the best alignment:



- ▶ A dynamic-programming search (Viterbi-like) for the highest-likelihood **monotonic** text→frame map.
- ▶ **Monotonic** = speech moves left→right, can't skip or repeat tokens → no "babbling".
- ▶ It also gives the **duration target** (frames per token) for the duration loss.



How training works

Per batch: text→prior · real audio→latent z · flow · **MAS** aligns text↔frames · slice a segment → decoder → waveform → losses. Then alternate two AdamW optimizers:

- ▶ **D step:** real vs. *detached* fake → discriminator loss
- ▶ **G step:** 5 losses: recon (L1 mel, **x45**) · KL · duration · adversarial · feature-matching

*I wrapped the alternation in **decorators** — read the method and it's linear; need the wiring, read the decorator. Logic in the function, complexity in the decorator.*



Abstraction

solving the readability problem — or creating it :)

The GAN loop is messy: two optimizers, detach, zero-grad, backward, clip, AMP scaler, accumulation, checkpointing, all wrapped around two tiny core steps.

- ▶ I moved **all that wiring into decorators**, leaving only the real **D-step / G-step logic** in the methods.
- ▶ Read the method → the algorithm is linear and clear. Need the plumbing → read the decorator.

Normal in ML: `torch.nn.Module` hides backward/autograd the same way. The complexity doesn't vanish, it goes where it belongs.



Where & how we trained

- ▶ **Colab first:** sessions dropped, free tier ~15 min, ephemeral disk. Painful.
- ▶ **Then local, RTX 5060 (8 GB).** No disconnects, full control.
- ▶ **Fitting 8 GB:** batch 1 + **gradient accumulation**, **AMP fp16**, `num_workers=0` (Windows DataLoader crash).

The 8 GB card is the whole reason for batch 1 + accumulation and fp16, full numbers & config coming up.



Problems & fixes: nothing converged at first

Early runs: recon stuck ~ 0.45 , duration loss **100–475**, the voice came out far too fast. Three bugs:

#	Problem	Fix
1	duration loss $\sim 100\times$ too large (averaged wrong)	<code>.sum()/mask.sum()</code> $\rightarrow$ dur 300 $\rightarrow$ 1.7
2	grad-clip too tight $\rightarrow$ killed the generator update	raise the clip $\rightarrow$ recon broke its plateau
3	fresh discriminator corrupted the pretrained decoder	warm up D 1500 steps + lower LR

$\rightarrow$ **real training: recon 0.85 $\rightarrow$ 0.40, correct speech rhythm.**



Training setup & cost


epoch	duration
1	1.89 h
2	1.89 h
3	1.96 h
4	2.02 h
5	1.93 h
total	= 9.68 GPU-h

local RTX 5060 (8 GB)

parameter	value
learning_rate	1e-4 vs 2e-4 from scratch
batch_size	1 + grad accumulation
segment_size	8192 decoder crop
grad_clip_norm	10
disc_warmup	1500 steps
use_amp	True fp16
loss_weights	mel 45 · kl 1 · adv 1 · fm 1


Plus **~3x more** in failed runs & debugging, and the **RVC baking** of the whole 13,100-clip dataset on top. The 9.68 h is only the final clean pass, the real work was everything before it.

Reconstruction loss (L1 mel)



Sharp drop, then plateau ~0.42, best at step **26,700 (0.404)**: the core voice is learned early; after that the GAN slowly drifts up. **That's why I don't take the last checkpoint.**


Duration loss



After the normalization fix, duration sits at a healthy **~1.6** (best 1.55, was 100–475): correct, stable speech rhythm.




KL divergence (posterior || prior)



Settles around **~1.7**: the text-prior successfully learns to match the audio posterior. Alignment works.




Adversarial loss (generator side)



Oscillates around ~ 4 without collapsing: a healthy GAN equilibrium, generator and discriminator stay balanced.


Discriminator loss



Stable around **-5.5**: neither winning nor collapsing. The warm-up kept it from overpowering the pretrained decoder.




Feature-matching loss



Early spike during warm-up, then settles **-2**: matches discriminator features, stabilizes the GAN.

Generator loss (total)



Total G loss ~28, dominated by recon×45. Flat & stable. **But a converging loss doesn't prove good speech** → that's what evaluation checks next.



Evaluation: what we measure & why

All metrics **lower** = better:

- ▶ **WER / CER** (Whisper): *intelligibility*, works for any voice, droid included. "Can a listener understand it?"
- ▶ **MCD**: *spectral* closeness to the **B1 droid reference**, the key one here. "Did it learn the droid timbre?"
- ▶ **The ceiling**: Whisper on the **real B1 recordings** = 5.75% WER → no model trained on this data can score below that.

The B1 twist: we score MCD against the **droid reference**, not the original female LJSpeech, otherwise we'd just be measuring "how un-droid-like" we are. A deliberate voice change only makes sense measured against the *new* target. (brief asks WER/CER + one spectral metric.)



How we evaluate

- ▶ **Held-out set:** the last 5% of the data, never seen in training.
- ▶ **Before / after:** one checkpoint swap, same prompts → apples-to-apples.
- ▶ **Contract-first eval module:** a manifest (text + reference audio) in → an `EvalResult` out; WER/CER/MCD in one pass.
- ▶ **Independent ASR:** Whisper is frozen and off-the-shelf, never part of training → no circularity.

```
manifest (text + reference audio) → synthesize @ checkpoint → Whisper ASR + spectral compare → WER + CER + MCD → EvalResult
```




Results: vs the ceiling

eval set	WER %	CER %
real B1 recordings (ceiling)	5.75	2.35
fine-tuned — epoch 2 (best)	14.39	7.01
fine-tuned — epoch 5	16.56	7.75

655 held-out clips (full 5%) - WER/CER vs text - Whisper - lower = better

Spectral / timbre: MCD to the droid reference **10.5** → **6.0 dB** → the voice moved to the droid.



same sentence: the harmonic structure shifts toward the droid timbre


Reading it: MCD **10.5** → **6.0 dB** = the timbre moved to the droid (fine-tuning worked). And intelligibility is close to the ceiling: our best model is **14.4%** WER vs the **5.75%** of the real B1 recordings, the model can't beat its teacher, but it lands near it.

Demo


Same sentence, before and after:

"The mission is complete. All systems are nominal. Returning to base."

1 · **Base VITS** (the pretrained voice, before)



2 · **Fine-tuned** → **droid** (same line, now the droid voice)



▶ same words, the **voice** changed.



Failure modes & limitations

An honest accounting of where it still falls short:

- ▶ **Intelligibility cost:** WER sits **above the ceiling** (~14% vs 5.75%) — understandable, but less clean than the base.
- ▶ **Timbre ≠ hard robot:** RVC moves a *speaker's timbre*, not the ring-mod buzz — that mechanical buzz would be a separate DSP effect.
- ▶ **GAN artifacts** early in training, and quality **over-fits** if pushed past the best epoch (why we don't take the last checkpoint).
- ▶ **Russian accent:** the original pre-processing idea, left as **future work**.



Takeaways & future work

- ▶ We **built & published a droid dataset** (RVC) and **fine-tuned VITS** onto that voice.
- ▶ We **re-implemented the GAN training** from scratch, discriminator, losses, the loop, on an inference-only model.
- ▶ Result: **the timbre moved to the droid** (MCD 10.5→6.0); intelligibility **~14% WER**, close to the **5.75% ceiling** of the data itself, a measured, bounded trade-off.
- ▶ **Future:** cut the WER by **freezing the text/timing front-end** (keeps durations stable), + the Russian accent.

Thank you

Questions & discussion

Dima, model & fine-tuning · Emir, Infra, data, engineering · Ilya, evaluation

github.com/Dmitri1S9/DL · 🧑‍🔬 [Dmitri13/ijlspeech-b1](#)